

## Exercise: 6

Meeting on 9th/ 11th February

### Aufgabe 1: Properties of the softmax function

The softmax function

$$\text{softmax}(\vec{x})_i = \frac{e^{x_i}}{\sum_j e^{x_j}} \quad (1)$$

is often used as the final activation function in a multi-class classification problem. Show/explain for the following properties of the softmax function:

- (a)  $\text{softmax}(\vec{x})_i$  corresponds to a probability distribution. Hence, show that  $\sum_i \text{softmax}(\vec{x})_i = 1$ .
- (b) For  $i$  it is for all  $j \neq i$ :  $x_i \gg x_j$ , then  $\text{softmax}(\vec{x})_i \rightarrow 1$  and  $\text{softmax}(\vec{x})_j \rightarrow 0$ . Hence, why is the function called *softmax*?


#### Solution:

$$(a) \sum_i \text{softmax}(\vec{x})_i = \sum_i \frac{e^{x_i}}{\sum_j e^{x_j}} = \frac{\sum_i e^{x_i}}{\sum_j e^{x_j}} = 1$$


- (b)  $x_i \gg x_j$  implies  $e^{x_i} \gg e^{x_j}$ , i.e.  $\sum_j e^{x_j} \rightarrow e^{x_i}$ , hence  $\text{softmax}(\vec{x})_i = 1$  and everything else 0. The maximum would be choosing only the largest value with probability/confidence of 1. Softmax allows for a degree of fuzziness. The training error is then not just 1 or 0, but something in between.

### Aufgabe 2: Backpropagation in a Deep Network - analytically

Consider a one-dimensional neural network with  $\tanh(x)$  as activation function. Forward-Pass and Backpropagation can be represented by:



Forward-Pass



Backpropagation

The loss function as well as forward and backward variables are computed as

$$\begin{aligned} L(f_3) &= \frac{1}{2}(y - f_3)^2 & r_0 &= h_0 \\ f_3(f_2) &= b_3 + w_3 f_2 & r_1 &= h_1 r_0 \\ f_2(f_1) &= \tanh(b_2 + w_2 f_1) & r_2 &= h_2 r_1 \\ f_1(x) &= \tanh(b_1 + w_1 x) & r_3 &= h_3 r_2 \end{aligned}$$

with

$$\begin{aligned} h_0 &= \frac{\partial L(f_3)}{\partial f_3} \\ h_1 &= \frac{\partial f_3(f_2)}{\partial f_2} \\ h_2 &= \frac{\partial f_2(f_1)}{\partial f_1} \\ h_3 &= \frac{\partial f_1(x)}{\partial x} \end{aligned}$$

Compute the derivatives in the variables  $h_i$  explicitly. You will need the following derivatives for backpropagation:

$$\begin{aligned} \frac{\partial L}{\partial w_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial w_3} \\ \frac{\partial L}{\partial b_3} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial b_3} \\ \frac{\partial L}{\partial w_2} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial f_2} \frac{\partial f_2}{\partial w_2} \\ \frac{\partial L}{\partial b_2} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial f_2} \frac{\partial f_2}{\partial b_2} \\ \frac{\partial L}{\partial w_1} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial f_2} \frac{\partial f_2}{\partial f_1} \frac{\partial f_1}{\partial w_1} \\ \frac{\partial L}{\partial b_1} &= \frac{\partial L}{\partial f_3} \frac{\partial f_3}{\partial f_2} \frac{\partial f_2}{\partial f_1} \frac{\partial f_1}{\partial b_1} \end{aligned}$$

Rewrite these equations using the forward variables  $f_i$  and the backward variables  $r_i$ .

**Solution:** The variables  $h_i$  are computed as:

- $h_0 = \frac{\partial L(f_3)}{\partial f_3} = \frac{\partial}{\partial f_3} \frac{1}{2}(y - f_3)^2 = -(y - f_3)$
- $h_1 = \frac{\partial f_3(f_2)}{\partial f_2} = \frac{\partial}{\partial f_2} (b_3 + w_3 f_2) = w_3$
- $h_2 = \frac{\partial f_2(f_1)}{\partial f_1} = \frac{\partial}{\partial f_1} \tanh(b_2 + w_2 f_1) = (1 - f_2^2) \cdot w_2$
- $h_3 = \frac{\partial f_1(x)}{\partial x} = \frac{\partial}{\partial x} \tanh(b_1 + w_1 x) = (1 - f_1^2) \cdot w_1$

The derivatives of the loss functions are:

$$\frac{\partial L}{\partial w_3} = h_0 f_2 = r_0 f_2$$

$$\frac{\partial L}{\partial b_3} = h_0 = r_0$$

$$\frac{\partial L}{\partial w_2} = h_0 h_1 (1 - f_2^2) f_1 = r_1 (1 - f_2^2) f_1$$

$$\frac{\partial L}{\partial b_2} = h_0 h_1 (1 - f_2^2) = r_1 (1 - f_2^2)$$

$$\frac{\partial L}{\partial w_1} = h_0 h_1 h_2 (1 - f_1^2) x = r_2 (1 - f_1^2) x$$

$$\frac{\partial L}{\partial b_1} = h_0 h_1 h_2 (1 - f_1^2) = r_2 (1 - f_1^2)$$

### Aufgabe 3: Convolution and Pooling operations in a CNN

Consider the following three matrices:

$$\bullet A = \begin{pmatrix} 4 & -4 & 1 & 3 \\ -5 & 1 & -1 & 0 \\ 6 & 4 & 7 & 4 \\ 18 & 3 & -5 & 11 \end{pmatrix} \quad \bullet B = \begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 8 & 7 & 6 & 5 \\ 4 & 3 & 2 & 1 \end{pmatrix}$$

As well as the filter matrix  $W$ :

$$W = \begin{pmatrix} -3 & 3 \\ -1 & 2 \end{pmatrix}$$

Using  $W$  perform the convolution operations on  $A$  and  $B$ , once without padding (= "valid") and once with padding while keeping the original resolution (= "same").

#### Solution:

- Without padding (Padding = "valid"):

$$A * W = \begin{pmatrix} 4 & -4 & 1 & 3 \\ -5 & 1 & -1 & 0 \\ 6 & 4 & 7 & 4 \\ 18 & 3 & -5 & 11 \end{pmatrix} * \begin{pmatrix} -3 & 3 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} -17 & 12 & 7 \\ 20 & 4 & 4 \\ -18 & -4 & 18 \end{pmatrix}$$

$$B * W = \begin{pmatrix} 1 & 2 & 3 & 4 \\ 5 & 6 & 7 & 8 \\ 8 & 7 & 6 & 5 \\ 4 & 3 & 2 & 1 \end{pmatrix} * \begin{pmatrix} -3 & 3 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} 10 & 11 & 12 \\ 9 & 8 & 7 \\ -1 & -2 & -3 \end{pmatrix}$$

- With padding = "same":

$$A * W = \begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 4 & -4 & 1 & 3 \\ 0 & -5 & 1 & -1 & 0 \\ 0 & 6 & 4 & 7 & 4 \\ 0 & 18 & 3 & -5 & 11 \end{pmatrix} * \begin{pmatrix} -3 & 3 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} 8 & -12 & 6 & 5 \\ 2 & -17 & 12 & 7 \\ -3 & 20 & 4 & 4 \\ 54 & -18 & -4 & 18 \end{pmatrix}$$

$$B * W = \begin{pmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 1 & 2 & 3 & 4 \\ 0 & 5 & 6 & 7 & 8 \\ 0 & 8 & 7 & 6 & 5 \\ 0 & 4 & 3 & 2 & 1 \end{pmatrix} * \begin{pmatrix} -3 & 3 \\ -1 & 2 \end{pmatrix} = \begin{pmatrix} 2 & 3 & 4 & 5 \\ 13 & 10 & 11 & 12 \\ 31 & 9 & 8 & 7 \\ 32 & -1 & -2 & -3 \end{pmatrix}$$

### Aufgabe 4: Programming a simple Neural Network

The lecture showed how you can program and train a simple neural network without any major frameworks. In this exercise you should try this yourself.

- Write a class "neuralNetwork" which initializes all necessary parameters, like input and output nodes as well as weights, and possesses the methods "train" and "query". Train and evaluate the network on the MNIST<sup>1</sup> dataset.
- Try using your own or downloaded images as inputs.

#### Solution:

##### (a) Easy Neural Network in Python:

```
1 import numpy as np
2 from scipy.special import expit
3 import csv
4
5
6 class neuralNetwork:
7     # Klasse, um ein fully connected neuronales Netz mit einer versteckten Schicht zu
8     #   ↳ erstellen
9     def __init__(self,
10                  inputnodes: int,
11                  hiddennodes: int,
12                  outputnodes: int,
13                  learningrate: float = 0.01):
14         self.inodes = inputnodes
15         self.hnodes = hiddennodes
16         self.onodes = outputnodes
17
18         # Elemente der Gewichtsmatrizen sind W_ih und W_ho
19         # Gewicht w_ij verbindet den Knoten i der vorherigen Schicht mit dem Knoten j
20         #   ↳ in der nächsten
21         # Zufällige Initialisierung der Gewichte ueber eine Normalverteilung (np.random
22         #   ↳ .normal)
```

<sup>1</sup><https://pjreddie.com/projects/mnist-in-csv/>

```

20     self.wih = np.random.normal(0.0, pow(self.hnodes, -0.5), (self.hnodes, self.inodes
    → ))
21     self.who = np.random.normal(0.0, pow(self.onodes, -0.5), (self.onodes, self.
    → hnodes))
22
23     self.lr = learningrate
24
25     # Aktivierungsfunktion ist sigmoid (auch expit genannt) fuer alle Schichten
26     self.activation_function = lambda x: expit(x)
27
28     # Fuehrt ein Gradient Descent Update mit gegebenem Batch aus
29     def train(self, inputs_list, targets_list):
30         # Eingabe (Bild + Ground Truth) zu 2D Array konvertieren
31         # Bilder sind bereits flattened, also ein Bild = eine Zeile
32         # und ein one-hot Label-Vektor = eine Zeile, damit zwei Matrizen
33         inputs = np.array(inputs_list, ndmin=2).T
34         targets = np.array(targets_list, ndmin=2).T
35
36         # Berechnung der versteckten Schicht
37         hidden_inputs = np.dot(self.wih, inputs)
38         # Anwendung der Aktivierungsfunktion
39         hidden_outputs = self.activation_function(hidden_inputs)
40
41         # Berechnung der Output-Schicht
42         final_inputs = np.dot(self.who, hidden_outputs)
43         # Anwendung der Aktivierungsfunktion
44         final_outputs = self.activation_function(final_inputs)
45
46         # Berechnung der Fehler der Knoten
47         # Fehler der Ausgabeschicht ist einfache Differenz, berechnet aus der
    → Lossfunktion (Berechnung im Uebungsblatt)
48         output_errors = targets - final_outputs
49         # Fehler der versteckten Schicht propagieren von oben
50         hidden_errors = np.dot(self.who.T, output_errors)
51
52         # Gewichtsupdate
53         # Exakte Berechnung: elementweise Ableitung der Lossfunktion
54         # Mit Backpropagation nur noch abhaengig von den umgebenden Schichten (
    → genaue Formel in der Loesung)
55
56         # Gewichte zwischen versteckt und Output
57         self.who += self.lr * np.dot((output_errors * final_outputs * (1.0 - final_outputs))
    → ,
58                                     np.transpose(hidden_outputs))
59
60         # Gewichte zwischen versteckt und Input
61         self.wih += self.lr * np.dot((hidden_errors * hidden_outputs * (1.0 -
    → hidden_outputs)), np.transpose(inputs))
62
63     # Berechnet einen Forward-Pass des Netzes
64     def query(self, inputs_list):
65         # Eingabe zu 2D Array konvertieren
66         inputs = np.array(inputs_list, ndmin=2).T
67
68         # Berechnung der versteckten Schicht

```

```

69     hidden_inputs = np.dot(self.wih, inputs)
70     # Anwendung der Aktivierungsfunktion
71     hidden_outputs = self.activation_function(hidden_inputs)
72
73     # Berechnung der Output-Schicht
74     final_inputs = np.dot(self.who, hidden_outputs)
75     # Finale Ausgabe berechnen
76     final_outputs = self.activation_function(final_inputs)
77
78     return final_outputs
79
80
81 # Anzahl der Knoten pro Schicht
82 input_nodes = 784 # aufgerollte MNIST-Auflösung von 28 x 28
83 hidden_nodes = 200 # Anzahl versteckter Knoten wählbar
84 output_nodes = 10 # 10 mögliche Klassen als Ausgabe
85
86 # Lernrate
87 learning_rate = 0.1
88
89 # Netzwerk initialisieren
90 n = neuralNetwork(input_nodes, hidden_nodes, output_nodes, learning_rate)
91
92
93 # MNIST-Daten als numpy array laden
94 # Aufbau der Daten: Jede Reihe enthält das Label als ersten Eintrag und aufgerolltes
    ↳ Bild als Rest
95 with open("mnist_train.csv", 'r') as train_file:
96     reader = csv.reader(train_file)
97     train_data = np.asarray(list(reader), dtype=int)
98 with open("mnist_test.csv", 'r') as test_file:
99     reader = csv.reader(test_file)
100    test_data = np.asarray(list(reader), dtype=int)
101
102
103 # Training
104
105 # Anzahl Epochen, also wie oft wird durch den Datensatz iteriert
106 epochs = 5
107
108 for e in range(epochs):
109     # Datensatz mischen
110     np.random.shuffle(train_data)
111     # Durch die Reihen iterieren
112     for row in train_data:
113         label = row[0]
114         inputs = row[1:]
115         # Inputs auf 0 bis 1 skalieren
116         inputs = (inputs / 255.0 * 0.999) + 0.001
117         # Label One-Hot encoden, also alles auf 0 bis auf Index = Label, den auf 1 (mit
            ↳ label smoothing)
118         targets = np.zeros(output_nodes) + 0.001
119         targets[label] = 0.999
120         n.train(inputs, targets)
121

```

```

122
123 # Testen
124
125 # Tracker, der die richtig und falsch erkannten Faelle speichert
126 scorecard = []
127
128 # durch den Test-Datensatz iterieren
129 for row in test_data:
130     # gleiche Aufbereitung wie beim Training
131     correct_label = row[0]
132     inputs = row[1:]
133     inputs = (inputs / 255.0 * 0.999) + 0.001
134     # Forward-pass aufrufen
135     outputs = n.query(inputs)
136     # Label ueber maximale Wahrscheinlichkeit bestimmen
137     label = np.argmax(outputs)
138     # Korrekte oder falsche Erkennung in die Evaluation eintragen
139     if (label == correct_label):
140         # 1 bei korrekter Erkennung
141         scorecard.append(1)
142     else:
143         # 0 bei falscher Erkennung
144         scorecard.append(0)
145
146 # Evaluation
147
148 # Performance des Netzes wird ueber die Accuracy bewertet, also korrekt erkannte
149     ↳ Bilder durch Anzahl aller Bilder
150 accuracy = np.sum(scorecard) / len(scorecard)
151 print("Accuracy = ", accuracy)

```

The code from the lecture has been used for the most part. Below some explanations on backpropagation.

Optimizing a network is always performed by minimizing the loss function  $L$ , which can theoretically be chosen arbitrarily. In this example we use the  $l_2$  loss, i.e. given prediction  $\vec{y}$  and targets  $\vec{t}$ :

$$L = \frac{1}{2} \sum_i^N (\vec{y} - \vec{t})^2$$

The weight update  $\Delta w_{jk}$  for a matrix element of the weight matrix  $W$  is computed via the derivative of the loss function (with the chain rule) and can be simplified using the backpropagation algorithm to:

$$\Delta w_{jk} = \alpha \cdot \delta(a_k) \cdot \frac{\partial}{\partial \hat{a}_k} \text{activation}(\hat{a}_k) \cdot a_j$$

where  $\delta(a_k)$  is the error of node  $a_k$  in the next layer,  $\hat{a}_k$  is the value of  $a_k$  before activation and  $a_j$  the value of node  $a_j$  in the previous layer. Since we use the sigmoid function everywhere and  $\sigma'(x) = \sigma(x)(1 - \sigma(x))$  and  $\sigma(\hat{a}_k) = a_k$ , the expression simplifies to

$$\Delta w_{jk} = \alpha \cdot \delta(a_k) \cdot a_k(1 - a_k) \cdot a_j.$$

We extend this formula to matrix multiplication in the code so we don't have to compute each element on its own (lines 56 to 61 in the code).

The error of the hidden layer is computed by summing all errors of each directly connected nodes in the next layer, weighted with the corresponding element of the weight matrix connecting the two nodes (line 50). The error of the output layer is in this case the difference between target and output (line 48), which is given by the derivative of the loss function.

(b) **Code fragment for reading in images:**

```

1 import imageio
2 import glob
3 import numpy as np
4
5 # Liste fuer unsere Daten
6 our_images = []
7 our_labels = []
8
9 for image_file_name in glob.glob('custom/meine_?.png'):
10     print ("loading ... ", image_file_name)
11     # Label wird hier ueber den Dateinamen bestimmt (auch andere Varianten moeglich
12     #   -> )
13     label = int(image_file_name[-5])
14     # als Graustufenbild reinladen
15     img_array = imageio.imread(image_file_name, as_gray=True)
16     # Von 28x28 zu 784 aufrollen
17     img_data = 255.0 - img_array.reshape(784)
18     # Passende Skalierung vornehmen
19     img_data = (img_data / 255.0 * 0.999) + 0.001
20     # Zur Liste mit allen Bildern und Labels hinzufuegen
21     our_images.append(img_data)
22     our_labels.append(label)
23
24 output = n.query(our_images)
25 results = [np.argmax(row) for row in output]
26
27 # print oder return (results, our_labels)

```

It is best to run it all in one script and save the network weights in a fitting format, to load them later without having to retrain the network. Jupyter notebooks are well suited for this task.

## Aufgabe 5: Q-Learning

Consider the *Markov Decision Problem* with three states  $Z \in \{1, 2, 3\}$  and their rewards  $B$  with  $B_1 = -1$ ,  $B_2 = -2$  and  $B_3 = 0$ . State  $Z = 3$  is a terminal state. In the states  $Z = 1$  and  $Z = 2$  there are two possible actions  $a$  and  $b$ . The indeterministic transition model is shown in Fig. 1.

In this scenario the agent does not know the transition model nor the rewards. Instead he know only about the number of states and the available actions.

- (a) Discuss, which action would be the most sensible one (as an external observer






Figure 1: The indeterministic transition model

with access to all information) to end with the highest possible reward (or the lowest penalty) in the terminal state.

**Solution:** The agent should try to reach  $Z = 3$  as quickly as possible since staying in one of the other two states only gives negative rewards. In state  $Z = 1$  they try action  $b$  to reach  $Z = 3$ . In state  $Z = 2$  they try action  $a$  to reach  $Z = 1$  and from there reach  $Z = 3$  with action  $b$  since there is less penalty for staying in 1.

(b) Explain the components of the formula  $Q(s, a)$  for the Q-learning algorithm.

**Solution:**

$$Q_{t+1}(s_t, a_t) = Q_t(s_t, a_t) + \alpha_t(s_t, a_t) \cdot \left( R_{t+1} + \gamma \cdot \max_a Q_t(s_{t+1}, a) - Q_t(s_t, a_t) \right)$$

- $Q_{t+1}(s_t, a_t)$ : New Q-Value
- $Q_t(s_t, a_t)$ : Old Q-Value
- $\alpha_t(s_t, a_t)$ : learning rate
- $R_{t+1}$ : reward
- $\gamma$ : Discount factor
- $\max_a Q_t(s_{t+1}, a)$ : Estimation of the future optimal value
- $R_{t+1} + \gamma \cdot \max_a Q_t(s_{t+1}, a)$ : learned value

Also see: English wikipedia page for Q-Learning!

(c) What is the meaning if the discount factor  $\gamma$ ?

**Solution:**  $\gamma$  weights the rewards by time, i.e. the smaller  $\gamma$  is the faster the agent tries to get rewards since the later rewards are smaller due to weighting.

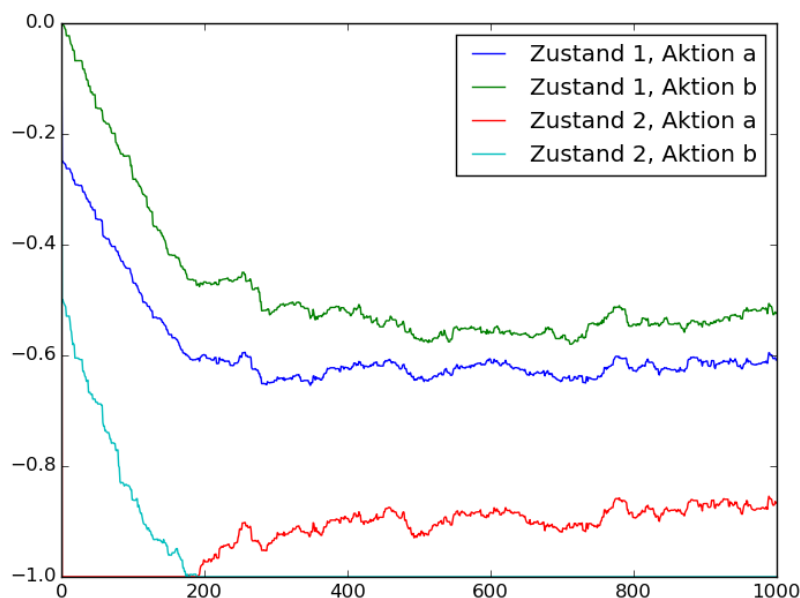
(d) Provide an outline of the sequence of the Q-learning algorithm.

**Solution:**

1. Initialize  $Q(s,a)$ -Matrix
2. For every episode do: Choose a random starting point and repeat until a terminal state is reached
  - Randomly choose a possible action in the current state  $s_t$
  - With that explore  $s_{t+1}$  (provide the agent with an unknown environment)
  - Compute  $Q_{t+1}(s_t, a_t)$  and with that update the  $Q$ -Matrix
  - Increment  $t$  (go from  $s_t$  to  $s_{t+1}$ )

(e) Program the Q-learning algorithm for  $\alpha = 0.001$  and show the evolution of the (normalized)  $Q$ -matrix across 1000 iterations.

**Solution:** Here is an exemplary parameter evolution. The max. value is normalized to  $-1$ :



You can see that the agent first considers action  $b$  to be better in state 2 but changes their opinion over time.

**Example code in Python:**

```

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 # world transition matrix (unknown to agent) T(s, a, s')
5 wtm = np.zeros((3, 2, 3))
6 wtm[0, 0, 0] = 0.2
7 wtm[0, 0, 1] = 0.8
8 wtm[0, 1, 0] = 0.9
9 wtm[0, 1, 2] = 0.1
10
11 wtm[1, 0, 0] = 0.8
12 wtm[1, 0, 1] = 0.2
13 wtm[1, 1, 1] = 0.9
14 wtm[1, 1, 2] = 0.1
15
16 reward = [-1, -2, 0]
17
18 n_episodes = 1000
19
20
21 def select_random_action():
22     # random 0 or 1
23     return np.random.randint(0, 2)
24
25
26 def get_new_state(action, cur_state):
27     r = np.random.rand()
28
29     t = 0
30     for i in range(3):
31         t += wtm[cur_state, action, i]
32         if r < t:
33             return i
34
35     raise Exception('Unexpected error', 'Maybe unallowed state and action')
36
37
38 def get_random_state():
39     # random 0, 1, or 2
40     return np.random.randint(0, 3)
41
42
43 def is_goal_state(state):
44     return state == 2
45
46
47 def compute_Q_value(Q, R, state, action, new_state, alpha=0.001, gamma=0.1):
48     return Q[state, action] + alpha * (R[state] + gamma * np.max(Q[new_state, :]) - Q[
49         ↪ state, action])
50
51 # ALGORITHM
52 # initialize Q-Matrix
53 Q = np.zeros((3, 2))
54 Qmemory = np.zeros((n_episodes, 3, 2))

```

```
55
56 for episode in range(n_episodes):
57     print('Episode %d' % episode)
58     state = get_random_state()
59
60     while not is_goal_state(state):
61         action = select_random_action()
62         new_state = get_new_state(action, state)
63
64         Q[state, action] = compute_Q_value(Q, reward, state, action, new_state)
65
66         state = new_state
67
68     # normalize Q (optional)
69     m = np.max(np.abs(Q))
70     if m > 0:
71         Q = Q / m
72
73     Qmemory[episode] = Q
74
75     # print('Result Q=%s' % Q)
76
77 plt.plot(Qmemory[:, 0, 0], label='Zustand 1, Aktion a')
78 plt.plot(Qmemory[:, 0, 1], label='Zustand 1, Aktion b')
79 plt.plot(Qmemory[:, 1, 0], label='Zustand 2, Aktion a')
80 plt.plot(Qmemory[:, 1, 1], label='Zustand 2, Aktion b')
81 plt.legend()
82 plt.show()
```