

I Artificial Intelligence

**II Problem Solving**

3. Solving Problems by Searching

**4. Search in Complex Environments**

5. Adversarial Search and Games

6. Constraint Satisfaction Problems

III Knowledge, Reasoning, Planning

IV Uncertain Knowledge and Reasoning

V Machine Learning

VI Communicating, Perceiving, and Acting

VII Conclusions



- Local Search and Optimization Problems
- Local Search in Continuous Spaces
- Search with Nondeterministic Agents
- Search in Partially Observable Environments
- Online Search Agents and Unknown Environments



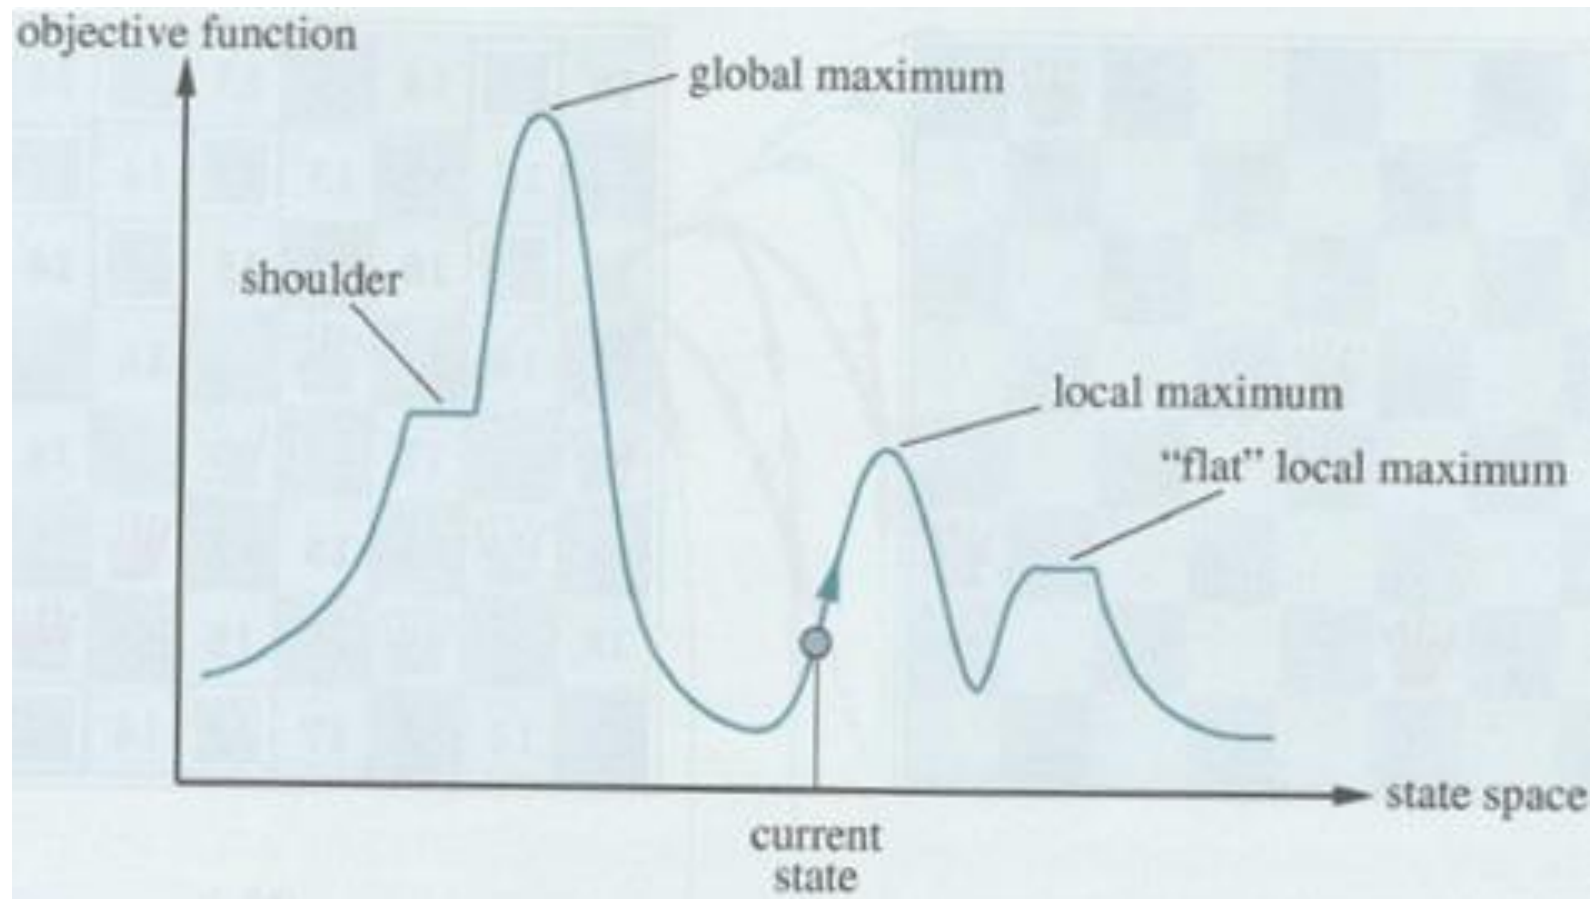
Frank Puppe

- Local search algorithms are useful, if only the solution is relevant, not the path to the solution.
  - Example: 8-queens-problem
  - Other Examples: VLSI-Layout, factory floor layout, job shop scheduling, automatic programming, telecommunication network optimization, crop planning, portfolio managment
- Idea: Start with a **complete-state formulation** and **improve it by stepwise variations**
- Algorithms:
  - **Hill-Climbing**
  - **Simulated Annealing**
  - **Local Beam Search**
  - **Genetic Algorithms**











- Choose alway the highest-value sucessor node and return the current node, if no improvement is possible
  - No search tree, no backtracking → very efficient
  - Problems: Local maximas, plateaus, shoulders, flat local maximas


```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  current ← problem.INITIAL
  while true do
    neighbor ← a highest-valued successor state of current
    if VALUE(neighbor) ≤ VALUE(current) then return current
    current ← neighbor
```



Frank Puppe

- Left: In each step, a queen is moved to reduce the total number of conflicts (e.g. to 12)
- Right: Local minima, where hill-climbing gets stuck

18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14		13	16	13	16
	14	17	15		14	16	16
17		16	18	15		15	
18	14		15	15	14		16
14	14	13	17	12	14	12	18



- **Sideways move:** Limited; to pass shoulders but to avoid infinite loops in e.g. flat maxima
- **Stochastic hill-climbing:** Random selection of all improvements
  - **First-choice hill-climbing:** Generate successor-nodes and take the first improvement
- **Random-restart hill-climbing:** Repeat hill-climbing with randomly generated initial states
- **Simulated Annealing:** Allow worsening with a low probability
- **Local beam search:** Simultaneous search an several paths





- Idea: Allow „down-hill“ steps to overcome local maxima
  - Model: Gradually cooling of hot material (e.g. metallurgy), freezing of water etc.
- Problem: Control of down-hill steps
- Solution:
  - Random factor for choosing steps: if it improves the solution, choose it with 100%, if not, choose it with a probability depending on the degree of worsening
  - Terminate, by lowering the probability of down-hill steps continuously

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  current  $\leftarrow$  problem.INITIAL
  for  $t = 1$  to  $\infty$  do
     $T \leftarrow$  schedule( $t$ )
    if  $T = 0$  then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE(current) – VALUE(next)
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{-\Delta E/T}$ 
```

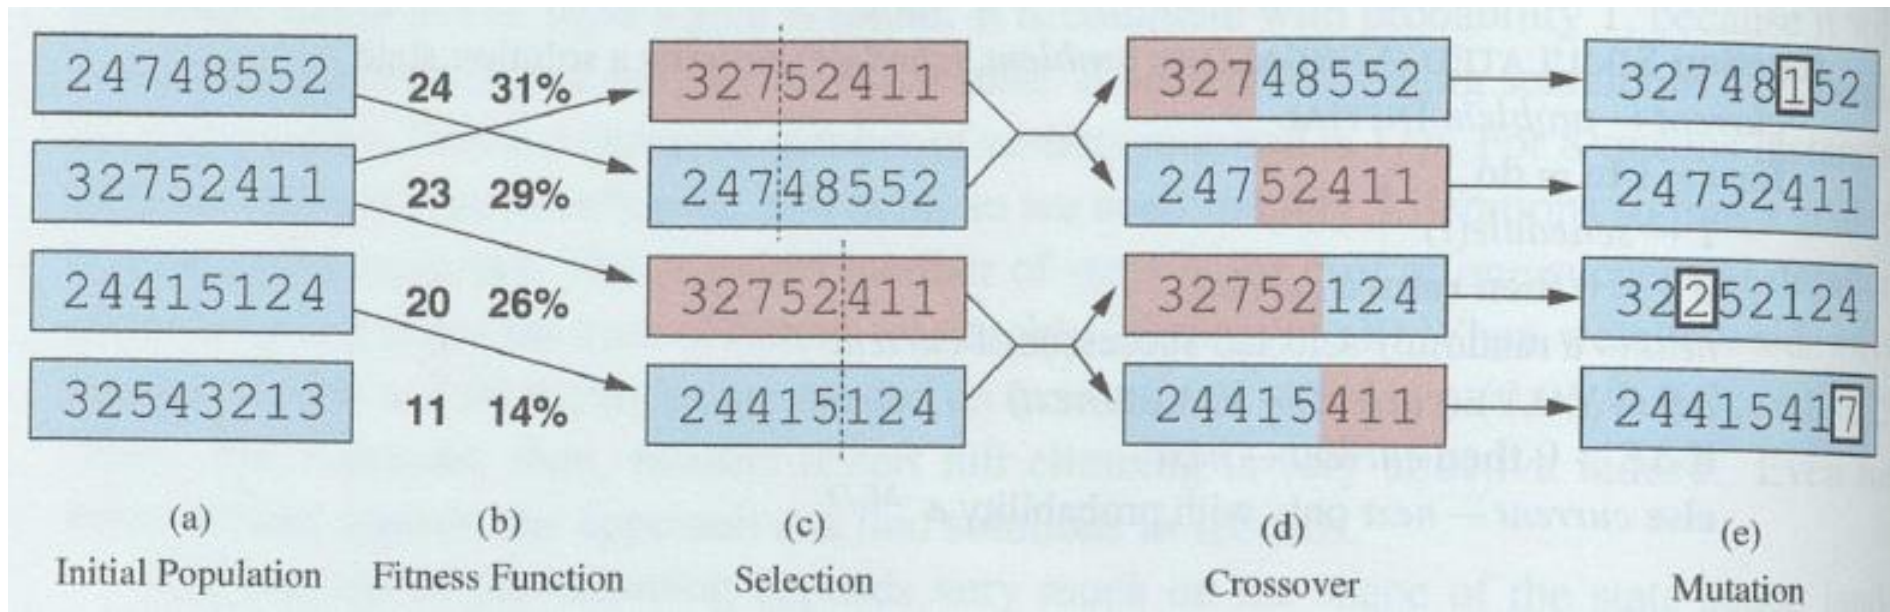


- **Idea:** Keep a set of  $k$  nodes („beam“) instead of one node. From one iteration to the next, the  $k$  best successors are chosen.
- **Difference to random restart hill-climbing:** Passing of information among parallel searches and concentration of search in promising regions.
- **Problem:** Concentration in a small region of the space (e.g. around a high, but local maxima).
- Improvement: **Stochastic generation of successors** nodes similar to stochastic hill-climbing for generation of diversity.





- Variant of stochastic beam search, which can **combine useful solution-parts** (blocks)
- Main steps (data structure is pool of solutions called „**population**“)
- Repeat until **termination criterium**
  1. **Select:** Select two solutions in population according to fitness
  2. **Recombine:** Generate new solution from both parents
  3. **Repair** (optionally): Repair (improve) new solution
  4. **Mutate:** Modify new solution randomly by mutations
- **Choose best solutions** (from new population)

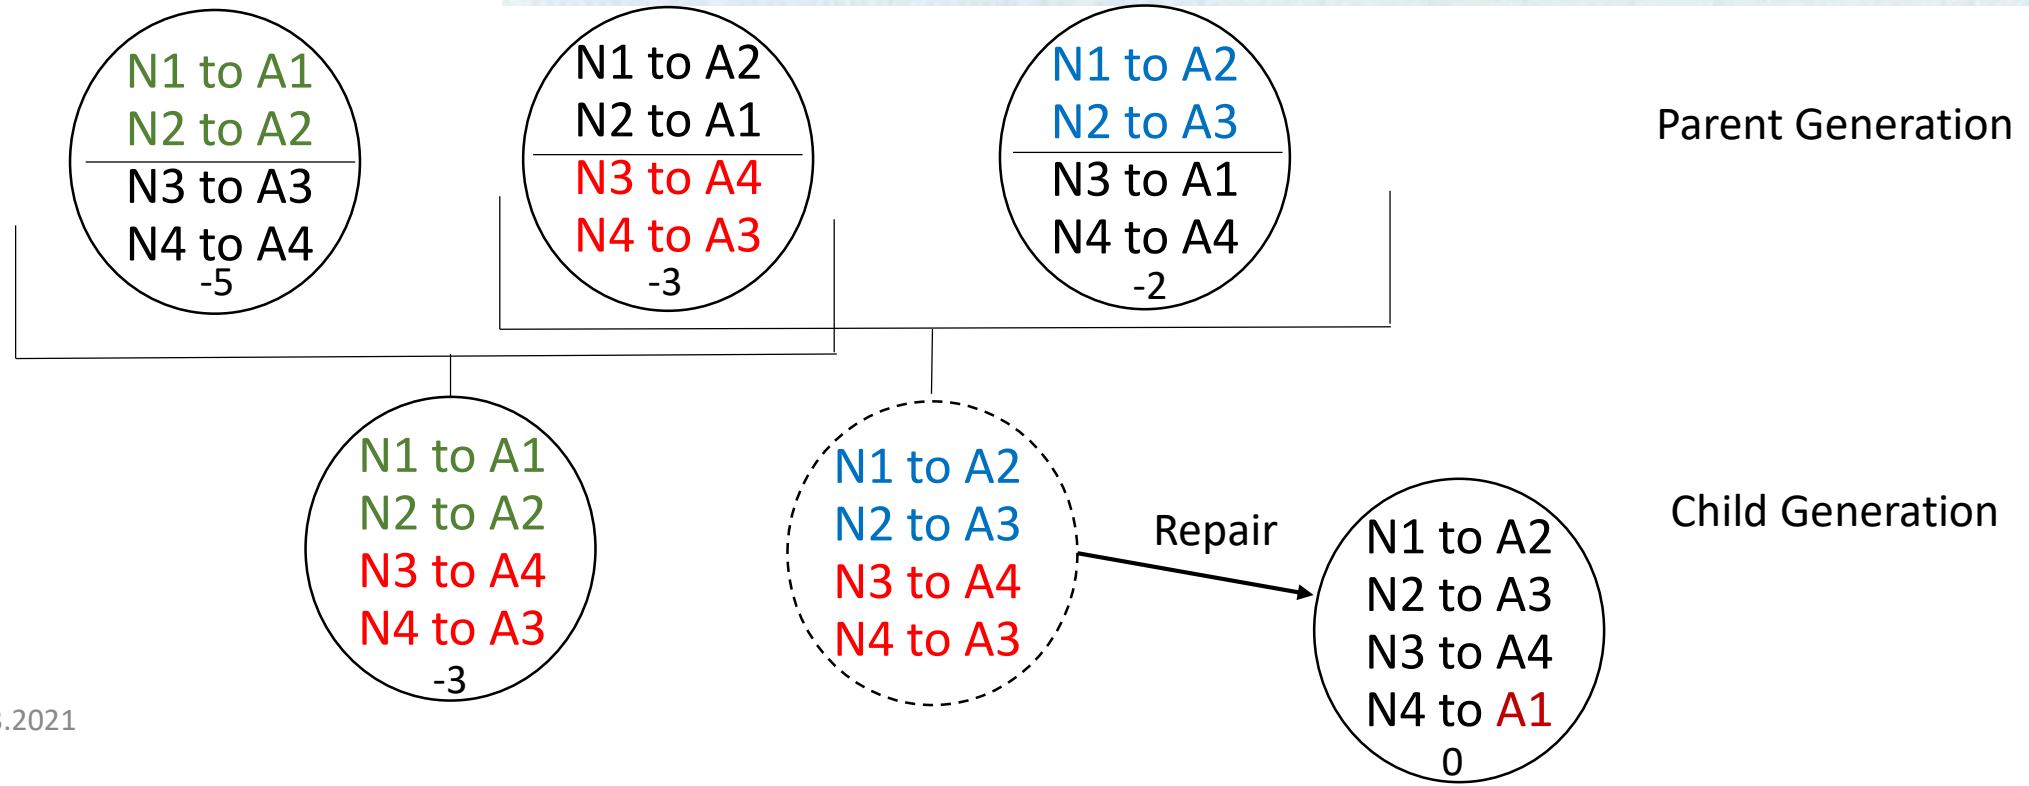



Frank Puppe

- **Size of population**
- **Representation of individuals**
  - Genetic algorithms: String over an alphabet (in biology DNA over alphabet ACGT)
  - Evolution strategies: Individual is a sequence of real numbers
  - Genetic programming: Individual is a computer program
- Mixing  $p$  parents:  $p = 1$ : Stochastic beam search;  $p = 2$ : Standard case;  $p > 2$ : Possible
- **Selection process**: Different functions usually depending on fitness
- **Recombination procedure**: Standard is to (randomly) select a crossover point
- **Mutation rate**: Determines, how often an offspring have random mutations
- **Makeup of next generation**: Keep top-scoring parents? Discard individuals below threshold?



Frank Puppe



```

function GENETIC-ALGORITHM(population, fitness) returns an individual
  repeat
    weights  $\leftarrow$  WEIGHTED-BY(population, fitness)
    population2  $\leftarrow$  empty list
    for i = 1 to SIZE(population) do
      parent1, parent2  $\leftarrow$  WEIGHTED-RANDOM-CHOICES(population, weights, 2)
      child  $\leftarrow$  REPRODUCE(parent1, parent2)
      if (small random probability) then child  $\leftarrow$  MUTATE(child)
      add child to population2
    population  $\leftarrow$  population2
  until some individual is fit enough, or enough time has elapsed
  return the best individual in population, according to fitness

function REPRODUCE(parent1, parent2) returns an individual
  n  $\leftarrow$  LENGTH(parent1)
  c  $\leftarrow$  random number from 1 to n
  return APPEND(SUBSTRING(parent1, 1, c), SUBSTRING(parent2, c + 1, n))
  
```

- **Population:**  
Ordered list of individuals
- **Weight:** List of fitness values for each individual
- **Fitness:**  
Function to compute weights





- Example problem: Place 3 airports in Romania with minimum square distance to the cities
  - Input: Coordinates of cities  $C_i$  (maybe with population weights), whose next airport is  $i$
  - Output: Coordinate  $(x,y)$  of the three airports  $(x_1, y_1; x_2, y_2; x_3, y_3)$
- Optimization criteria:

$$f(\mathbf{x}) = f(x_1, y_1, x_2, y_2, x_3, y_3) = \sum_{i=1}^3 \sum_{c \in C_i} (x_i - x_c)^2 + (y_i - y_c)^2$$



Frank Puppe

- **Discretization of neighborhoods**

- Moving an airport in x or y-direction with a constant d
- With 6 variables 12 successors per state
- Application of any search algorithms

- **Local gradient search**

- Derivation of goal function for each variable
- Gradient of goal function is a vector indicating the size and direction of the steepest slope ( $\alpha$  = step size)
- Newton-Raphson algorithm often effective, also for matrices


$$\frac{\partial f}{\partial x_1} = 2 \sum_{c \in C_1} (x_1 - x_c)$$

$$\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$$

- **Linear programming**

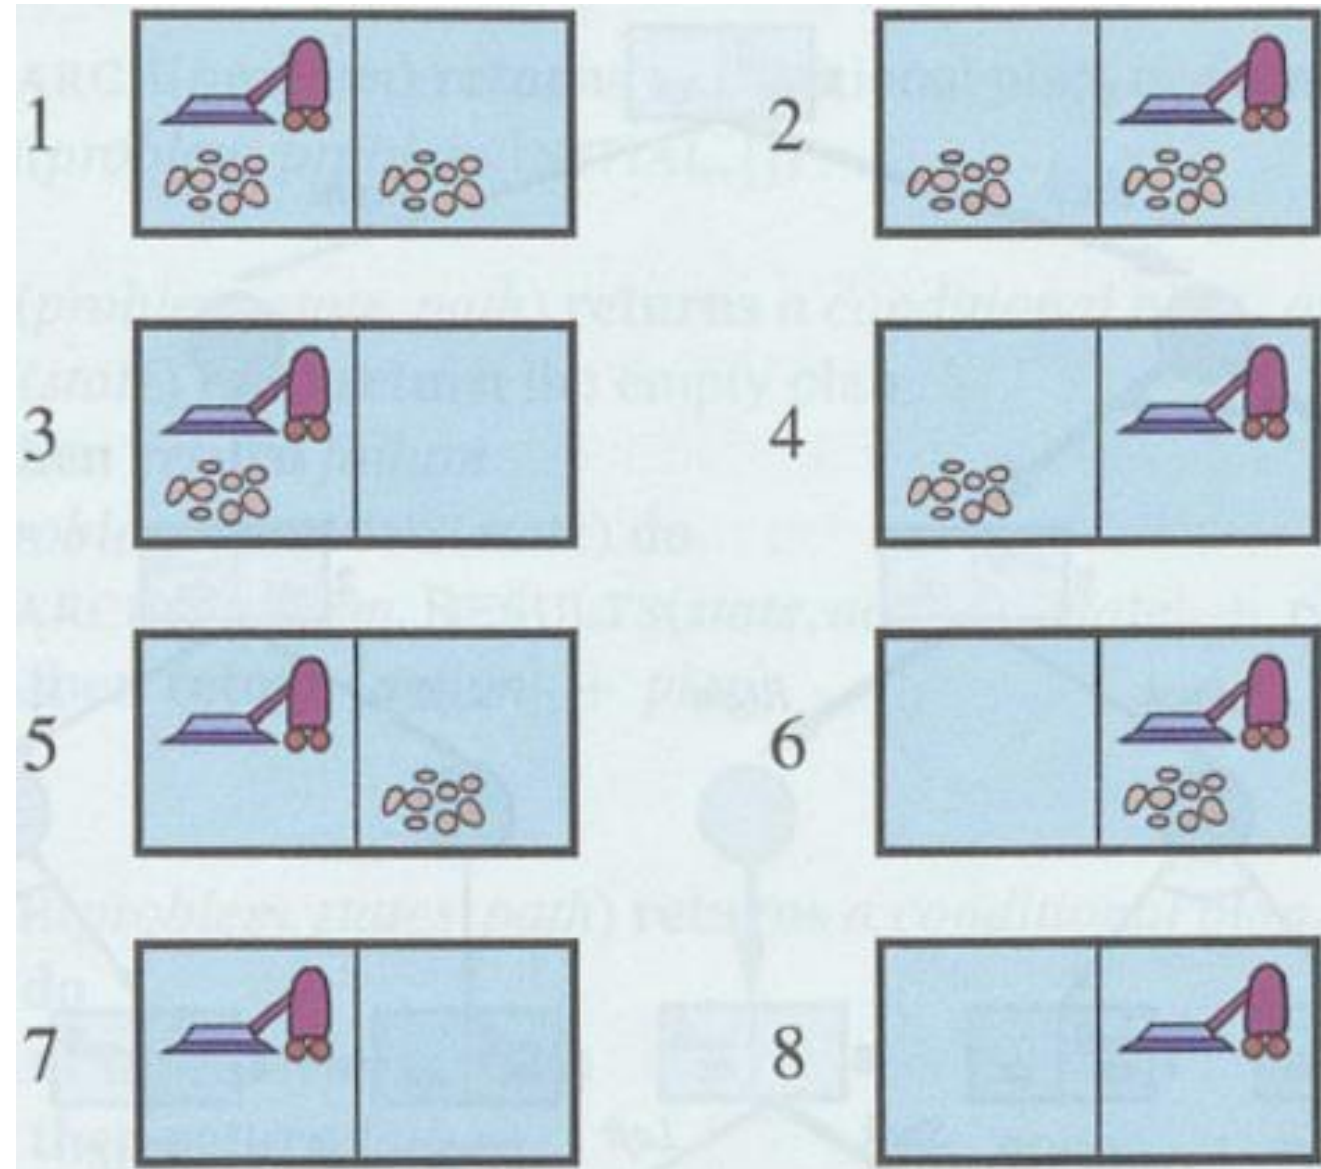
- Constraints for goal function must be linear (e.g. airports should not be placed in mountains)
- Goal function must be linear too (e.g. sum of distances)

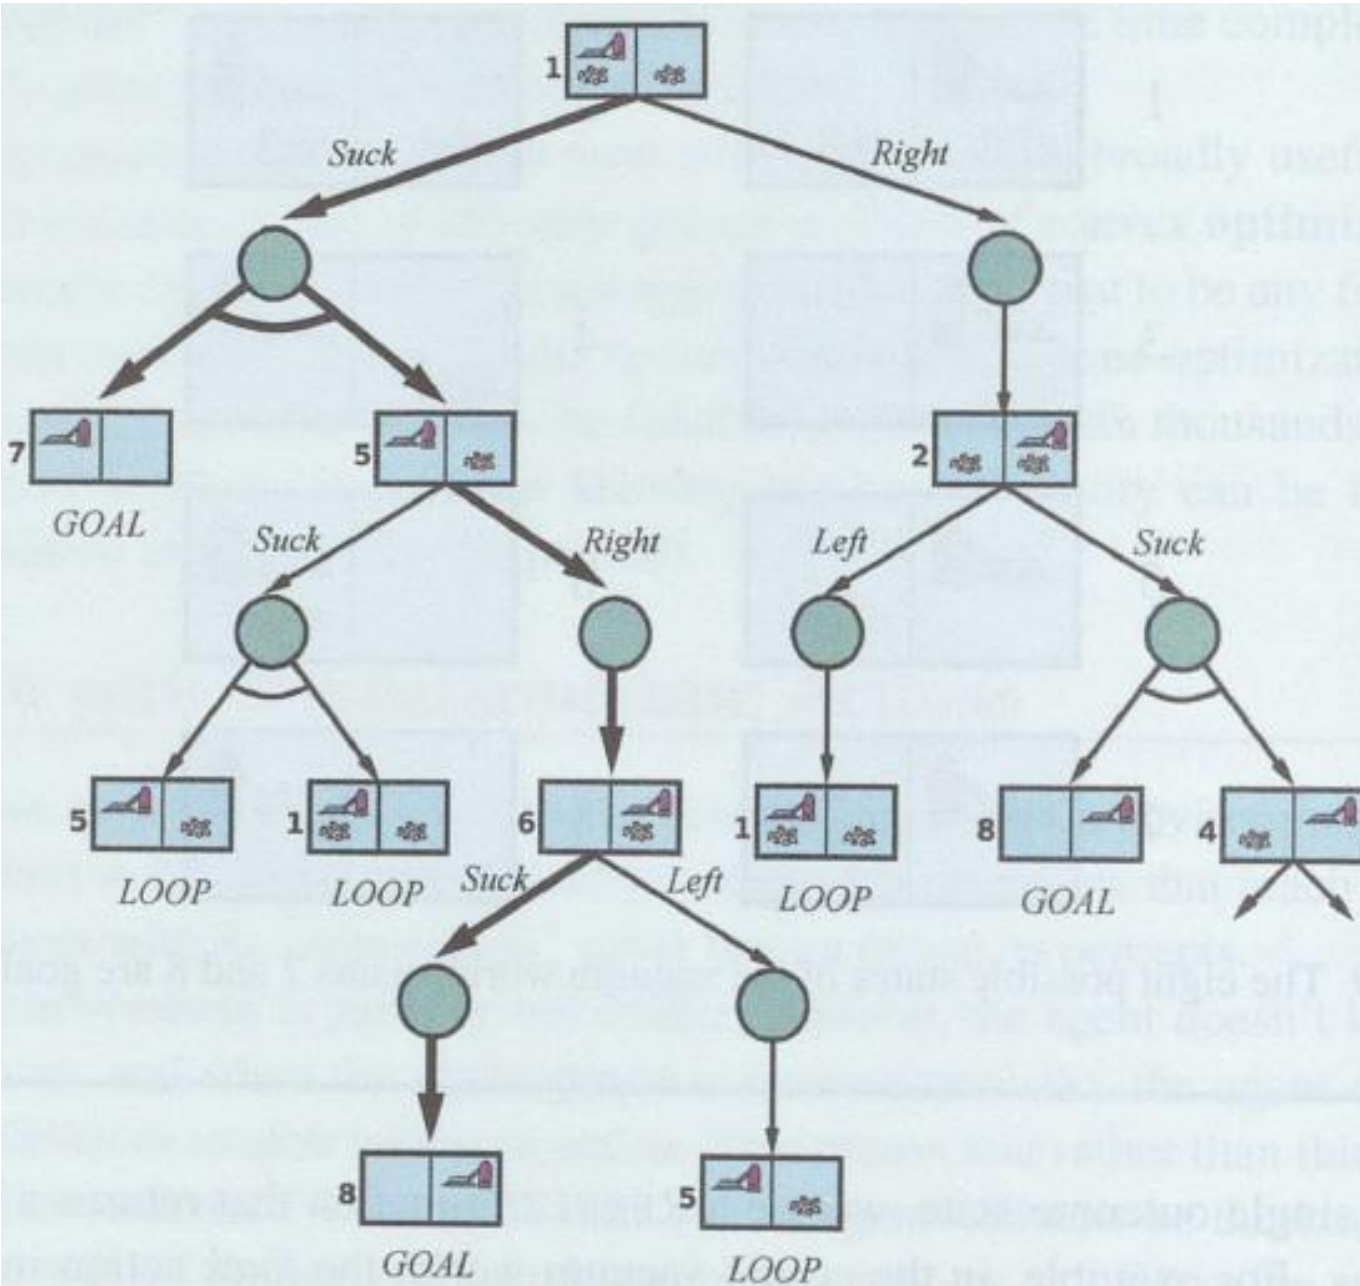
Newton-Raphson for solving equations of the form  $g(x) = 0$   
 $x \leftarrow x - g(x)/g'(x)$






- Example problem:
  - Vacuum world with erratic actions:
    - Suck of a dirty square makes it clean and sometimes cleans the adjacent square too
    - Suck of a clean square sometimes deposits dirt on it
- Solution in state 1
  - Deterministic: No plan available
  - Conditional plan:
    1. Suck
    2. if state = 5  
then [Right, Suck]  
else [ ]





- OR-Nodes: rectangles
- AND-Nodes: circles
  - All results must be handled
  - The state must be checked
- Solution is a subtree of the complete search tree that has ...
  - a goal node at every leaf
  - specifies one action for each of its OR-nodes
  - includes every outcome branch at each of its AND-nodes
- Bold arrows: Solution

- Different search algorithms possible, e.g. depth-first, breadth-first, best-first
- Next slide: Recursive depth-first algorithm
- Key aspect: dealing with cycles (arising often in nondeterministic problems)
  - Finding a plan avoiding cycles
  - Keep trying an indeterministic action until the desired outcome occurs
    - Example in a slippery vacuum world, where movements may fail, i.e. the agent may stay in the same location:
      - [Suck, while State = 5 do Right, Suck]
      - [Suck,  $L_1$ : Right, **if** State = unchanged **then**  $L_1$  **else** Suck]
- Might result in infinite loop
- Try a limited number of repetitions (like inserting an electronic card)





Solution is a conditional plan considering every nonterministic outcome

```

function AND-OR-SEARCH(problem) returns a conditional plan, or failure
    return OR-SEARCH(problem, problem.INITIAL, [])

function OR-SEARCH(problem, state, path) returns a conditional plan, or failure
    if problem.IS-GOAL(state) then return the empty plan
    if IS-CYCLE(path) then return failure
    for each action in problem.ACTIONS(state) do
        plan  $\leftarrow$  AND-SEARCH(problem, RESULTS(state, action), [state] + path)
        if plan  $\neq$  failure then return [action] + plan
    return failure



function AND-SEARCH(problem, states, path) returns a conditional plan, or failure
    for each  $s_i$  in states do
         $plan_i \leftarrow$  OR-SEARCH(problem,  $s_i$ , path)
        if  $plan_i =$  failure then return failure
    return [if  $s_1$  then  $plan_1$  else if  $s_2$  then  $plan_2$  else ... if  $s_{n-1}$  then  $plan_{n-1}$  else  $plan_n$ ]
    
```



- Example problem: Vacuum world without or with partial sensor information
- Solution approaches: **Search in belief states** and update of belief states
  - **Belief state:** All possible states compatible with the current information
    - Should include the physical state



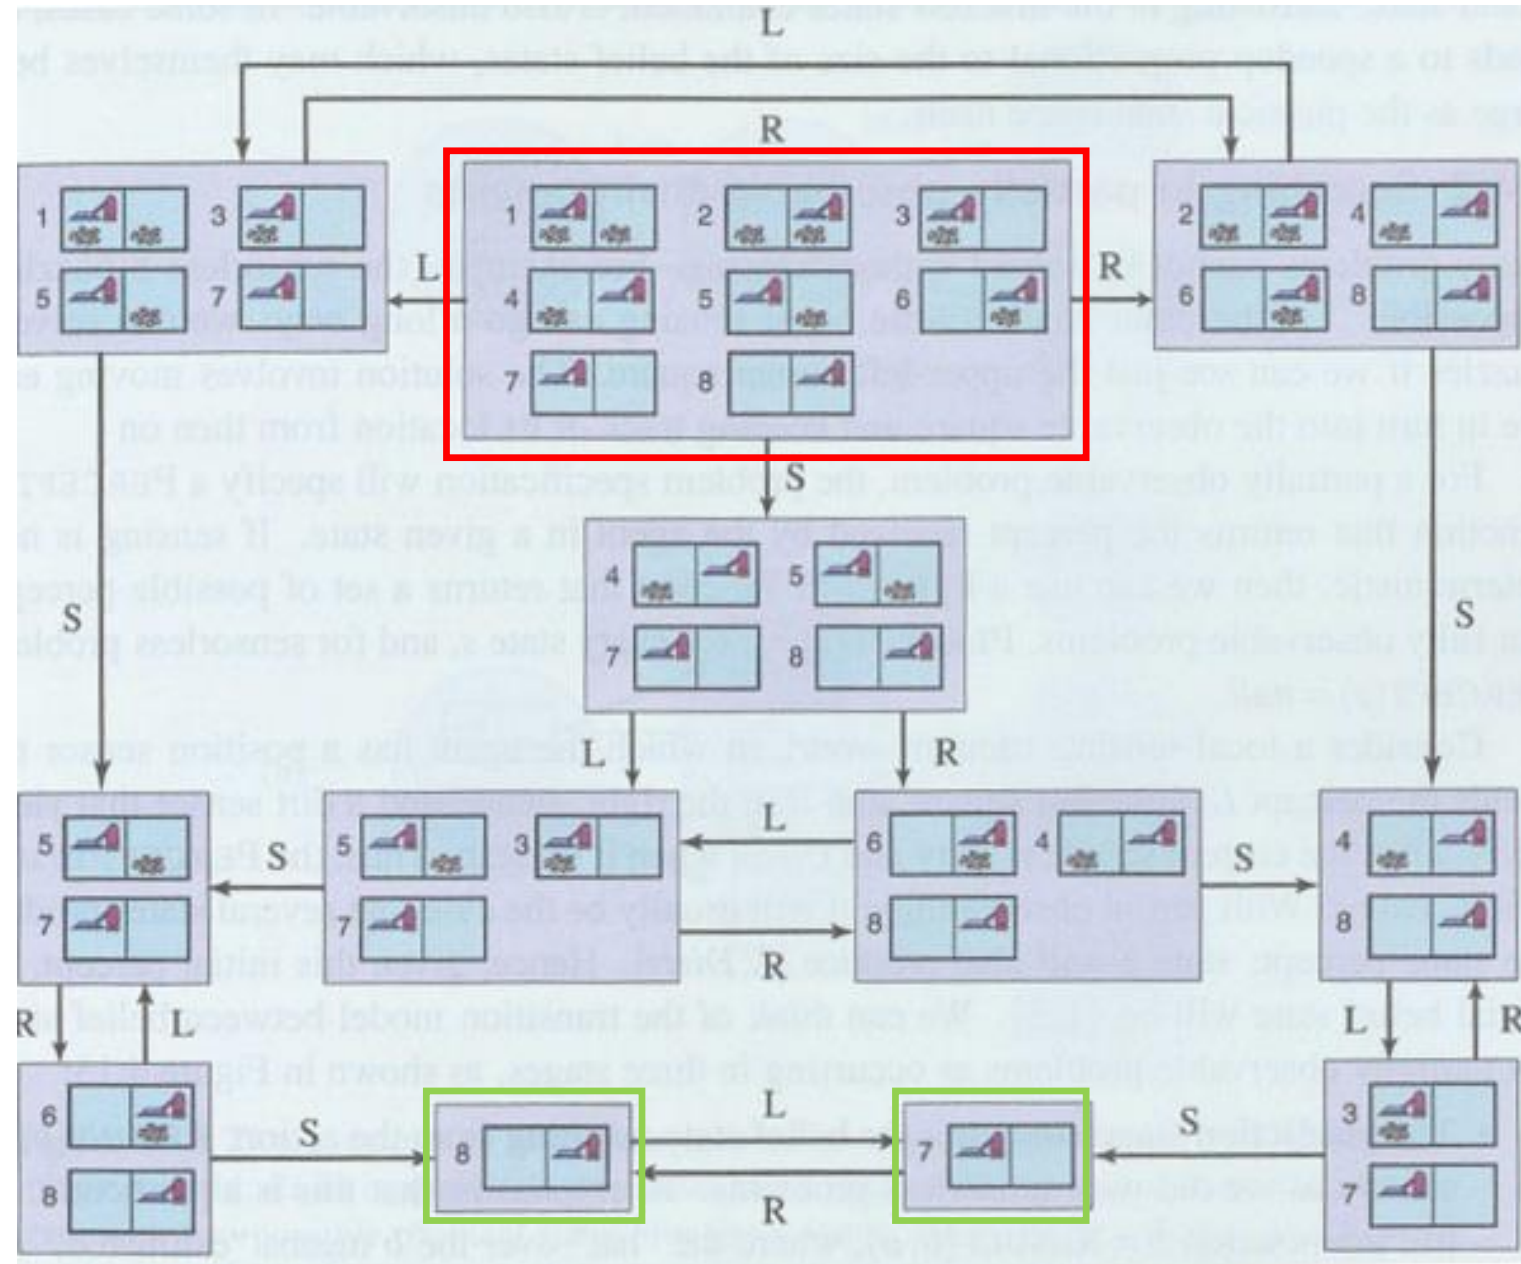
- Sensorless (conformant) problem
- Advantages: Sensors are expensive and often not reliable
- Examples:
  - Producing a base state for machines (e.g. restart a computer)
  - Infections: Prescribing broadband antibiotics
  - Sensorless vacuum world:
    - Effect of the deterministic action „right“ in a belief state:
    - Effect of the indeterministic action „right“ in a belief state:





... with deterministic actions

- Start state in red
- goal states in green
- Possible solutions:  
[right, suck, left, suck]  
[left, suck, right, suck]



- **Belief state:** Set of all possible physical states
  - With N physical states, there are  $2^N$  belief states
- **Initial state:** Without knowledge all  $2^N$  belief states
- **Actions:**
  - Assumption: Illegal actions have no effect in environment
    - New belief state  $b' =$  union of all actions a in every physical state p of the current belief state b
  - Assumption: Illegal actions might cause (great) damage
    - It is safer to allow only actions, which are legal in all the states (computed by intersection)

$$\text{ACTIONS}(b) = \bigcup_{s \in b} \text{ACTIONS}_P(s)$$



- **Transition model:**

- For deterministic actions: the new belief state  $b$  has one result state  $s$  for each of the possible actions  $a$ : set  $b' = \text{Result}(b, a) = \{s' : s' = \text{Result}_p(s, a) \text{ und } s \in b\}$
- For indeterministic action: the new belief state may have several result states for each of the possible actions: set  $b' = \text{Result}(b, a) = \{s' : s' \in \text{Result}_p(s, a) \text{ und } s \in b\} = \bigcup_{s \in b} \text{RESULTS}_p(s, a)$
- The size of the result set may decrease for deterministic actions and may increase for indeterministic actions.
- **Goal test:** All physical states  $p$  in belief state  $b$  should satisfy the goal test
- **Action costs:** Could be difficult, if the same action has different costs if applied to different states, otherwise straightforward.



- Apply **ordinary search algorithms** in belief state representation
  - Special treatment of supersets: if a superset  $\{1, 3, 5, 7\}$  is solved, all subsets, e.g. (e.g.  $\{5, 7\}$ ) are solved too.
  - The representation of one belief state is already exponential
    - Compacter representation of belief states by using logic instead of enumeration, e.g. after action [right]: „not in left cell“
  - Very large space  $2^N$  instead of  $N$ , and even  $N$  is often too large for e.g. breadth-first search or  $A^*$  search.
- **Incremental belief state search**
  - Search a solution for just one physical state in the belief state
  - Check, whether this plan works also for the other physical states
  - If not, search for another plan in the first physical state etc.
  - Is efficient to find out, whether a problem has no solution




- Sensorless planning often impossible (e.g. 8-puzzle)
- A little bit of sensing often useful (e.g. one square in 8-puzzle sufficient)
- Algorithm aspects
  - Prediction step similar to sensorless planning
    - But results in possible percepts, that could be observed in predicted belief state
  - Update step computes for each possible percept new belief state from percept
    - For deterministic sensing, belief states for different possible percepts are disjoint
  - AND-OR search algorithm applicable
    - Solution is a conditional plan; agent tests condition and execute appropriate branch
    - Agent updates its belief state after each action (and percept)
    - Probabilistic information for indeterministic actions would be useful (treated later)






- A robot has 4 sonar sensors telling him, whether there is a wall, e.g. 1011 means North: Wall; East: No wall; South and West: Wall.
- Robot has 4 action to move in each direction.
  - However, move is non-deterministically, so the robot may end on any adjacent field.



(a) Possible locations of robot after  $E_1 = 1011$



(b) Possible locations of robot after  $E_1 = 1011, E_2 = 1010$







- Offline-Search: Solution is computed before execution
- Online-Search: Computation and Execution overlap
- Necessary in unknown environments (but goal and action with costs are known)
- Problems: Irreversible action leading in dead ends (e.g. robot drops off a cliff, one-way streets)



Frank Puppe

- **Online-Depth-First Search:**

- In each state, follow the action given by depth-first search
- If no action possible, go back to last branch (must be stored) and try another action
- In maze (right), the agent would walk:  $(1,1) \rightarrow (1,2) \rightarrow (1,1) \rightarrow (2,1) \rightarrow (2,2) \rightarrow (2,3) \rightarrow (1,3) \rightarrow (2,3) \rightarrow (2,2) \rightarrow (2,1) \rightarrow (3,1) \rightarrow (3,2) \rightarrow (3,3)$
- Indefinite paths would prevent the agent from finding the goal (right below)
- Online variant of **iterative deepening depth-first search** would avoid indefinite paths and would find „shallow“ goals much faster.



- **Hill-climbing search** (does not explore the environment, might get stuck)
- Hill-climbing with random restarts and random walk (very slow)
- Augmenting hill-climbing with memory: For each state, an heuristic cost-estimate to reach the goal is stored and updated with more experience
- **Real-time A\* (LRTA\*)**: Agent builds a map in its result table:
  - The estimated cost to reach a goal through neighbor state  $s'$  is cost to get to  $s'$  + estimated cost to goal, i.e.:  $c(s, a, s') + H(s')$ .
  - Example: cost increase in red circle to escape plateau
  - Prefer untried states  $u$  by  $H(u) = 0$

