

- I Artificial Intelligence
- II Problem Solving
- III Knowledge, Reasoning, Planning
- IV Uncertain Knowledge and Reasoning**
 - 12. Quantifying Uncertainty
 - 13. Probabilistic Reasoning
 - 14. Probabilistic Reasoning over Time
 - 15. Probabilistic Programming
 - 16. Making Simple Decisions
 - 17. Making Complex Decisions**
 - 18. Multiagent Decision Making
- V Machine Learning
- VI Communicating, Perceiving, and Acting
- VII Conclusions



- Sequential Decision Problems
- Algorithms for Markov Decision Processes (MDPs)
- Bandit Problems
- Partially Observable MDPs (POMDPs)
- Algorithms for Solving POMDPs



Frank Puppe

- **Problem:**



- Often, an agent cannot reach the goal in one step, but needs multiple steps
- If the outcome of each action (step) is indeterministic, what is the optimal sequence?
- Indeterminism is represented by a transition model:
 - For each state and action, the probability of successor states is specified
 - Our known techniques (planning and replanning) ignore probabilities and utilities

- **Solution approach:**

- Specify for each state a ruleset (policy) depending on available evidence
 - The ruleset defines implicitly an exponential number of sequential plans
- The derivation of the rules depends on state utilities and action probabilities




- The agent wants to maximize its reward
 - For reaching a terminal state (4,2) and (4,3) the reward is -1 resp. +1.
 - For each movement (action) the reward is -0.04
- In each state, the agent can choose among four actions (up, down, left, right) and has a probability of 80% to reach the intended state, and 20% to move at right angles to the intended direction
 - Collision with a wall results in no movement
- If the agent reaches the state (4,3) e.g. in 10 steps, its total utility is $9 * -0.04 + 1 = +0.64$



- Specification:
 - Set of states (with an initial state s_0)
 - Set of actions in each state
 - Transition model $P(s' | s, a)$
 - Reward function $R(s, a, s')$
- Total utility depends on sequence of states and actions (environment history) and is the sum of the reward functions for reaching each state
- Solution is not a fixed action sequence (because of indeterminism), but a policy π
 - The policy recommends for each state s an action $\pi(s)$
- Optimal policy π^* yields the highest expected utility



- Left: Optimal policy with reward of -0.04 between nonterminal states
 - In state (3,1) there are two policies, because both „Up“ and „Left“ are optimal
- Right: Different policies for different rewards between nonterminal states



- The optimal policy also depends on the time horizon:
 - **Finite horizon:** Fixed time N , after which nothing matters (game over)
 - Optimal policy depends on how much time is left: **Nonstationary policy**
 - **Infinite horizon** without fixed deadline
 - **Stationary policy** not depending on time left (simpler)
 - But may have terminal states
 - Often with additive discounted rewards
 - Future rewards are less valuable than current rewards
 - Utility of a history: $U_h([s_0, a_0, s_1, a_1, s_2, \dots]) = R(S_0, a_0, s_1) + \gamma R(S_1, a_1, s_2) + \gamma^2 R(S_2, a_2, s_3) + \dots$
 - **Discount factor** γ : number between 0 and 1
 - Justification of discount factor: empirical (intuitive for humans and animals), economic (rewards can be re-invested), unsure future, ...
 - With discounted rewards, the utility of an infinite sequence is finite
 - **Proper policy:** Guaranteed to reach a terminal state (so we can use $\gamma = 1$)
 - Infinite sequences can be compared in terms of the **average reward** per time



- **Utility of an action sequence:** Sum of discounted rewards of the actions in each state
- **Expected utility of a policy:** Evidence (E) weighted sum of discounted rewards (R) of actions of policy $\pi(S_t)$

$$U^\pi(s) = E \left[\sum_{t=0}^{\infty} \gamma^t R(S_t, \pi(S_t), S_{t+1}) \right]$$
- **Optimal policy π_s^* :** Policy with highest expected utility from all policies:

$$\pi_s^* = \operatorname{argmax}_{\pi} U^\pi(s)$$
 - For infinite horizons the optimal policy π^* is independant from the start state, because it states for each state the same policy



Computing the optimal policy for a state is choosing the action leading to the neighbor states with highest utility:

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} [\sum_{s'} P(S' | s, a) [R(s, a, s') + \gamma U(s')]]$$

Computing the utility of a state depends on the utilities of the neighboring states:

$$U(s) = \max_{a \in A(s)} [\sum_{s'} P(s' | s, a) [R(s, a, s') + \gamma U(s')]]$$

Thus we get n equations with n unknowns, the **Bellman equations**

3	0.8516	0.9078	0.9578	+1
2	0.8016		0.7003	-1
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4

example utilities (with $\gamma = 1$; $R = -0.04$)



For simplicity with discount factor $\gamma = 1$

$U(1,1) =$

$$\begin{aligned} &\max \{ [0.8(-0.04 + U(1,2)) + 0.1 (-0.04 + U(2,1)) + 0.1 (-0.04 + U(1,1))], \\ &\quad [0.9(-0.04 + U(1,1)) + 0.1 (-0.04 + U(1,2))] \\ &\quad [0.9(-0.04 + U(1,1)) + 0.1 (-0.04 + U(2,1))] \\ &\quad [0.8(-0.04 + U(2,1)) + 0.1 (-0.04 + U(1,2)) + 0.1 (-0.04 + U(1,1))] \} \\ &= 0.7453 \end{aligned}$$

where the 4 expressions correspond to UP, Left, Down, Right

3	0.8516	0.9078	0.9578	+1
2	0.8016		0.7003	-1
1	0.7453	0.6953	0.6514	0.4279
	1	2	3	4



The Q-function $Q(s,a)$ is the expected utility for an action in a state. Thus:

$$U(s) = \max_a Q(s,a)$$

$$\pi^*(s) = \operatorname{argmax}_a Q(s,a)$$

The Bellman equation can be rewritten with Q-functions:

$$Q(s,a) = \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma U(s')] = \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma \max_{a'} Q(s',a')]$$

General function Q-value (mdp, s, a, U) **returns** a utility value

$$\text{return } \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma U[s']]$$



- Solving n equations with n unknowns, but the equations are nonlinear because of the max operator
- Iterative approach: Start with arbitrary initial values for each state, update the states and repeat this until reaching an equilibrium
 - Value iteration: Iterate, until the utilities of the states converge
 - Policy iteration: Iterate, until the policy stops changing
 - (Linear programming)
 - (Online algorithms for MDPs)



- Compute utilities of states iteratively until the biggest change (δ) is very small ($\leq \epsilon(1-\gamma)/\gamma$)
- Bellman update: $U_{i+1}(s) \leftarrow \max_{a \in A(s)} \sum_{s'} P(s'|s,a) [R(s,a,s') + \gamma U_i(s')]$ or
 $U_{i+1}(s) \leftarrow \max_{a \in A(s)} \text{Q-value}(\text{mdp}, s, a, U)$

```

function VALUE-ITERATION(mdp,  $\epsilon$ ) returns a utility function
  inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s'|s,a)$ ,
           rewards  $R(s,a,s')$ , discount  $\gamma$ 
            $\epsilon$ , the maximum error allowed in the utility of any state
  local variables:  $U, U'$ , vectors of utilities for states in  $S$ , initially zero
                      $\delta$ , the maximum relative change in the utility of any state

  repeat
     $U \leftarrow U'; \delta \leftarrow 0$ 
    for each state  $s$  in  $S$  do
       $U'[s] \leftarrow \max_{a \in A(s)} \text{Q-VALUE}(\text{mdp}, s, a, U)$ 
      if  $|U'[s] - U[s]| > \delta$  then  $\delta \leftarrow |U'[s] - U[s]|$ 
  until  $\delta \leq \epsilon(1-\gamma)/\gamma$ 
  return  $U$ 
  
```



Frank Puppe

Example for Value Iteration

$$U(3,3) = \max \{ \begin{aligned} &[0.8 * 1 + 0.1 (-0.04 + U(3,2)) + 0.1 (-0.04 + U(3,3))], && // \text{Right} \\ &[0.8(-0.04 + U(3,3)) + 0.1 * 1 + 0.1 (-0.04 + U(2,3))] && // \text{Up} \\ &[0.8(-0.04 + U(2,3)) + 0.1 (-0.04 + U(3,3)) + 0.1 (-0.04 + U(3,2))] && // \text{Left} \\ &[0.8(-0.04 + U(3,2)) + 0.1 * 1 + 0.1 (-0.04 + U(2,3))] \} && // \text{Down} \end{aligned}$$

3	0	0	0	1
2	0		0	-1
1	0	0	0	0
	1	2	3	4

Computations 1. Iteration:

$$U(3,3): \text{Right: } 0.8 * 1 + 0.1 * -0.04 + 0.1 * -0.04 = 0.792$$

$$U(3,2): \text{Left: } 0.8 * -0.04 + 0.1 * -0.04 + 0.1 * -0.04 = -0.04$$

... (other states = -0.04)

3			→	1
2			←	-1
1				↓
	1	2	3	4

3	-0.04	-0.04	0.792	1
2	-0.04		-0.04	-1
1	-0.04	-0.04	-0.04	-0.04
	1	2	3	4

Computations 2. Iteration:

$$U(3,3): \text{Right: } 0.8 * 1 + 0.1 * (-0.04 - 0.04) + 0.1 * (0.792 - 0.04) = 0.8672$$

$$U(3,3): \text{Left: } 0.8 * (-0.04 - 0.04) + 0.1 * (0.792 - 0.04) + 0.1 * (-0.04 - 0.04) = 0.0032$$

$$U(3,2): \text{Up: } 0.8 * (0.792 - 0.04) + 0.1 * -1 + 0.1 * (-0.04 - 0.04) = 0.4936$$

$$U(2,3): \text{Right: } 0.8 * (0.792 - 0.04) + 0.2 * (-0.04 - 0.04) = 0.5856$$


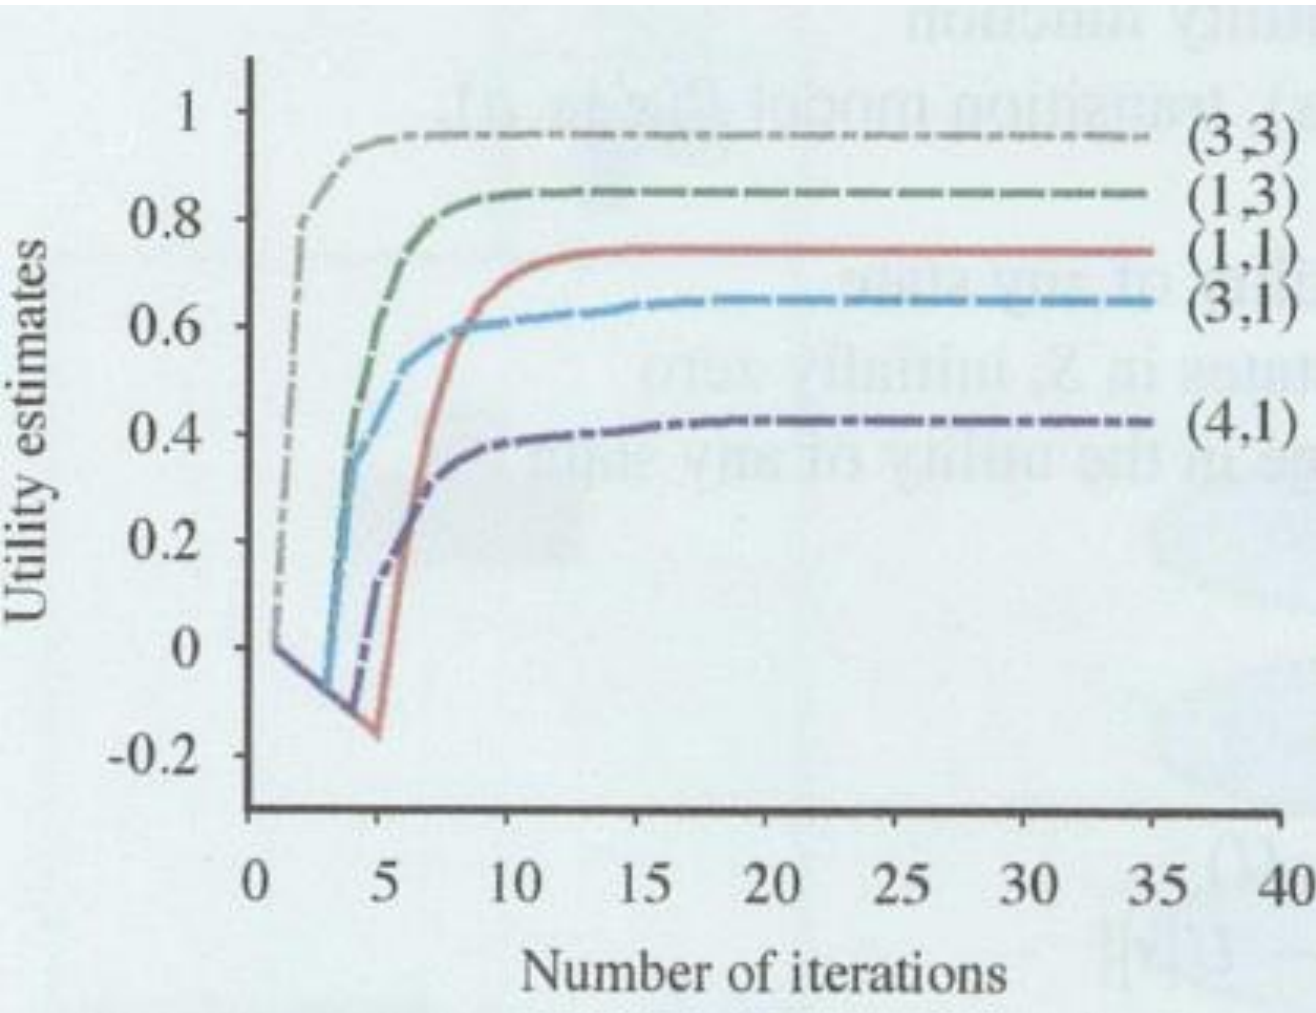
3		→	→	1
2			↑	-1
1				↓
	1	2	3	4

...

3	-0.08	0.5856	0.8672	1
2	-0.08		0.4936	-1
1	-0.08	-0.08	-0.08	-0.08
	1	2	3	4

Computations 3. Iteration:

...



The value iteration algorithm can be viewed als value propagation; the resulting policy stabelizes quickly



- What is the required exactness of values to find an optimal policy?
 - Sufficient exact to select the best policy.
- Iterate until the policy does not change from one iteration to the next



- Policy iteration algorithm alternates between two steps, beginning with some initial policy π
 - Policy evaluation: Given a policy π_i , calculate its utility U_i
 - Policy improvement: Calculate a new policy π_{i+1} based on U_i
- Termination, when the policy improvement step yields no change in the utilities


```
function POLICY-ITERATION(mdp) returns a policy
  inputs: mdp, an MDP with states  $S$ , actions  $A(s)$ , transition model  $P(s' | s, a)$ 
  local variables:  $U$ , a vector of utilities for states in  $S$ , initially zero
                    $\pi$ , a policy vector indexed by state, initially random

  repeat
     $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, \textit{mdp})$ 
     $\textit{unchanged?} \leftarrow \text{true}$ 
    for each state  $s$  in  $S$  do
       $a^* \leftarrow \underset{a \in A(s)}{\text{argmax}} \text{ Q-VALUE}(\textit{mdp}, s, a, U)$ 
      if  $\text{Q-VALUE}(\textit{mdp}, s, a^*, U) > \text{Q-VALUE}(\textit{mdp}, s, \pi[s], U)$  then
         $\pi[s] \leftarrow a^*$ ;  $\textit{unchanged?} \leftarrow \text{false}$ 
    until  $\textit{unchanged?}$ 
  return  $\pi$ 
```




Frank Puppe

- **Analytic version:** Solving n equations with n unknowns with linear algebra methods in $O(n^3)$, since the policy is known
 - Example with $\pi_i(1,1) = U_p$, $\pi_i(1,2) = U_p$, etc.
 - $U_i(1,1) = 0.8(-0.04 + U_i(1,2)) + 0.1(-0.04 + U_i(2,1)) + 0.1(-0.04 + U_i(1,1))$
 - $U_i(1,2) = 0.8(-0.04 + U_i(1,3)) + 0.2(-0.04 + U_i(1,2))$
 - ...
- **Iterative version:** Simplified version of Bellman update in value iteration with given policy
 - $U_{i+1}(s) \leftarrow \sum_{s'} P(s'|s, \pi_i(s)) [R(s, \pi_i(s), s') + \gamma U_i(s')] \quad // \text{ i.e. } Q\text{-value}(mdp, s, \pi_i(s), U)$
 - Must be repeated several times to produce next utility estimate
- **Asynchronous policy iteration:**
 - It is not necessary to update all states
 - Choosing any subset and perform either kind of updating of utilities is possible



- Value iteration and policy iteration are offline algorithms
 - Generate on optimal solution and then execute it
 - Not successful for large problems (like e.g. Tetris)
 - Approximate solutions are based on reinforcement learning
- Online algorithms are similar to game playing algorithms
 - Significant amount of computation at each decision point
 - Adaptation of ExpectiMiniMax algorithm for probabilistic games




Frank Puppe

- Difference to ExpectiMiniMax: Rewards on terminal and non-terminal nodes
- Evaluation function for non-terminal leaves of the tree
- With a low discount factor future rewards are less important allowing to cut the search tree
- Alternate to heuristic evaluation function: Generate N samples from probability distribution P and use the mean: $\sum_x P(x) f(x) \approx 1/N \sum_{i=1}^N f(x_i)$
- Improvements:
 - The search „tree“ is really a graph, since states appear several times
 - Constrain values of identical states to identical values
 - Adaptation of Bellman equations could be used
 - General approach: Real-time dynamic programming (RTDP)
 - Similar to learning real-time A* (LRTA*) algorithm
 - Useful for moderate sized domains with a high enough chance of repeated states
 - Apply reinforcement learning or an adaption of the Monte Carlo approach for games



- In Las Vegas, a **one-armed bandit** is a slot machine. A gambler can insert a coin, pull the lever and collect the winnings (if any). An **n-armed bandit** has n levers.
 - Behind each lever is a fixed, but unknown probability distribution of winnings
- General concept: Choose the best alternative (e.g. medical treatment, possible investment, project funding etc.) based on some samples
- Bandit problems can be defined as Markov reward processes





- The two bandits can be viewed as one one-armed-bandit with the reward of λ if not pulling an arm
- With just one arm, the only choice is to pull again or to stop (forever)
- Parameter T: Stopping time

With discount factor $\gamma = 0.5$ the bandits yield the followings utilities:

$$U(M) = 1.0 * 0 + 0.5 * 2 + 0.5^2 * 0 + 0.5^3 * 7.2 = 1.9$$

$$U(M_1) = \sum_{t=0}^{\infty} 0.5^t = 2.0$$

If switching is allowed once, the best option is switching after the fourth reward:

$$U(S) = 1.0 * 0 + 0.5 * 2 + 0.5^2 * 0 + 0.5^3 * 7.2 + \sum_{t=4}^{\infty} 0.5^t = 2.025$$



What is the value of λ of the one-armed bandit, so that an optimal strategy (with the best stopping time) is equivalent to stop immediately?

$$\max_{T>0} E \left[\left(\sum_{t=0}^{T-1} \gamma^t R_t \right) + \sum_{t=T}^{\infty} \gamma^t \lambda \right] = \sum_{t=0}^{\infty} \gamma^t \lambda$$

which simplifies to

Gittins index:

$$\lambda = \max_{T>0} \frac{E \left(\sum_{t=0}^{T-1} \gamma^t R_t \right)}{E \left(\sum_{t=0}^{T-1} \gamma^t \right)}$$

Computation of
Gittins index for
example:

T	1	2	3	4	5	6
R_t	0	2	0	7.2	0	0
$\sum \gamma^t R_t$	0.0	1.0	1.0	1.9	1.9	1.9
$\sum \gamma^t$	1.0	1.5	1.75	1.875	1.9375	1.9687
ratio	0.0	0.6667	0.5714	1.0133	0.9806	0.9651




Frank Puppe

- Optimal policy for any bandit problem:
 - Pull the arm with the highest Gittins index, then update the Gittins index
- Can be expanded into an equivalent MDP problem



Frank Puppe

- **Bernoulli Bandit:** Each arm M_i produces a reward of 0 (failure = f_i) or 1 (success = s_i) with a fixed, but unknown probability μ_i
 - State of an arm is defined by s_i and f_i and the transition probability predicts the next outcome to be 1 by the ratio $s_i / (s_i + f_i)$
 - Counts for s_i and f_i are initialized to 1, so that the initial probability is $\frac{1}{2}$
 - Example with probability p for getting a result R after each sample
 - Should balance selecting the arm with the highest payoff with exploring rarely used arms



Frank Puppe

- Calculating Gittins indices for more realistic problems is rarely easy
- Good approximations for combining high pay-off and exploration:
 - Algorithms based on **upper confidence bounds (UCB) heuristic**
 - Compute for each arm a confidence interval
 - Choose the arm with the highest upper bound of the confidence interval
 - $UCB(M_i) = \hat{\mu}_i + g(N)/\sqrt{N_i}$
 - Mit $\hat{\mu}_i$ = geschätzter_Mittelwert $/\sqrt{N_i}$ = proportional zur Standardabweichung;
 $g(N)$ = geeignet gewählte Funktion, z.B. $g(N) = (2\log(1 + N\log^2 N))^{1/2}$
 - Algorithm based on **Thompson sampling**:
 - Choose an arm randomly according to the probability that the arm is in fact optimal, given the samples so far
 - Simple implementation: Generate for each arm one sample based on its current probability of success and choose the arm that generated the best sample



- Example for bandit problem: Testing new medical treatments on seriously ill patients
 - Goal: Maximizing the total number of successes (saved lives) over time
- Alternate example scenario: Testing different drugs on samples of bacteria
 - Difference: No additional costs for test failures (like lost lives)
 - Implies a tendency to more exploration than bandit problems
 - Goal: Trying to make a good decision as fast as possible („**selection problem**“)
 - Candidate solution approach: Monte Carlo tree search algorithm for games with UCB selection heuristic
- Generalisation of bandit process: **Bandit superprocess**, where each arm is a full MDP
 - Several tasks, where one can attend to only one task at a time (e.g. selecting from multiple projects)
 - Selecting the locally best project not sufficient, but the time for delaying other projects is also important (e.g. spend more money than necessary for a project to be able to do quickly other attractive projects)
 - Computation of „opportunity costs“ necessary: How much utility is given up per time step by not devoting that time step to other projects?



- MDP: Environment fully observable
- POMDP: Environment partial observable, i.e. the agent does not know, in which state it is
 - POMDP has same elements as MDP and in addition a sensor model
 - Transition model $P(s' | s, a)$
 - Actions $A(s)$
 - Reward function $R(s, a, s')$
 - Sensor model $P(e, s)$, i.e. probability of perceiving evidence e in state s
 - Example: 4 x 3 world, where the agent does not know its location, but only has noisy four-bit sensor representing the presence or absence of a wall
 - From sensor model and evidence, agent computes belief state, i.e. a probability distribution of possible states
 - An action changes the belief state
 - POMDPs include the value of information (i.e. actions for improving the belief state)



- The value iteration algorithm for MDPs can be transferred to POMDPs requiring some extensions
 - Is hopelessly inefficient for larger problems – even for the 4 x 3 world.
- Online algorithms for POMDP are more straightforward
 - Start with some prior belief state
 - Choose an action based on some deliberation process on the current belief state
 - Update belief state based on new observation
- Good candidate: Expectimax algorithm for MDPs with belief states instead of physical states
 - Sampling at chance nodes to cut down branching factor
 - For large state space, approximate filtering algorithm like **particle filtering** („sequential importance sampling with resampling“ to focus on high probability regions of state space)
 - For problems with long horizons balance exploration and exploitation similar to MCTS (Monte Carlo Tree Search) with UCT selection policy using UCB (upper confidence bounds)
 - Successful only, if playout in MCTS has a chance of gaining positive rewards



- (right) Part of ExpectiMax tree with a uniform initial belief state and update after one action and appropriate noisy sensor information
- (below) An example sequence of percepts (error rate $\varepsilon = 20\%$), belief states and actions

